

The Coolest Way To Generate Combinations

Frank Ruskey* and Aaron Williams**

Dept. of Computer Science, University of Victoria.

Abstract. We present a practical and elegant method for generating all (s, t) -combinations (binary strings with s zeros and t ones): Identify the shortest prefix ending in 010 or 011 (or the entire string if no such prefix exists), and rotate it by one position to the right. This iterative rule gives an order to (s, t) -combinations that is circular and genlex. Moreover, the rotated portion of the string always contains at most four contiguous runs of 0s and 1s, so every iteration can be achieved by transposing at most two pairs of bits. This leads to an efficient loopless and branchless implementation that consists only of two variables and six assignment statements. The order also has a number of striking similarities to colex order, especially its recursive definition and ranking algorithm. In light of these similarities we have named our order *cool-lex*!

Keywords: Gray code order, combinations, binary strings, colex, loopless algorithm, branchless algorithm, prefix rotation, prefix shift

1 Background and Motivation

An important class of computational tasks is the listing of fundamental combinatorial structures such as permutations, combinations, trees, and so on. Regarding combinations, Donald E. Knuth writes in his upcoming volume of *The Art of Computer Programming* [8] “Even the apparently lowly topic of combination generation turns out to be surprisingly rich, I strongly believe in building up a firm foundation, so I have discussed this topic much more thoroughly than I will be able to do with material that is newer or less basic.”

The applications of combination generation are numerous and varied, and Gray codes for them are particularly valuable. We mention as application areas cryptography (where they have been implemented in hardware at NSA), genetic algorithms, software and hardware testing, statistical computation (e.g., for the bootstrap, Diaconis and Holmes [3]), and, of course, exhaustive combinatorial searches.

As is common, combinations are represented as binary strings, or bitstrings, of length $n = s + t$ containing s 0s and t 1s. We denote this set as $\mathbf{B}(s, t) = \{b_1b_2 \cdots b_n \mid \sum b_i = t\}$. Another way of representing combinations is as increasing sequences of the elements in the combination. We denote this set as $\mathbf{C}(s, t) = \{c_1c_2 \cdots c_t \mid 1 \leq c_1 < c_2 < \cdots < c_t \leq s + t\}$.

* Research supported in part by an NSERC Discovery Grant.

** Research supported in part by a NSERC PGS-D.

Our initial motivation was to consider the problem of listing the elements of $\mathbf{B}(s, t)$ so that successive bitstrings differ by a prefix that is cyclically shifted by one position to the right. We refer to such shifts as *prefix shifts*, or *rotations*, and they may be represented by a cyclic permutation $\sigma_k = (1\ 2\ \dots\ k)$ for some $2 \leq k \leq n$, where this permutation acts on the indices of a bitstring.

As far as we are aware, the only other class of strings that has a listing by prefix shifts are permutations, say of $\{1, 2, \dots, n\}$. In Corbett [1] and Jiang and Ruskey [7] it is shown that all permutations may be listed circularly by prefix shifts. That is, the directed Cayley graph with generators $(1\ 2)$, $(1\ 2\ 3)$, \dots , $(1\ 2\ \dots\ n)$ is Hamiltonian. In our case we have the same set of generators acting on the indices of the bitstring, but the underlying graph is not vertex-transitive; in fact, it is not regular.

There are many algorithms for generating combinations. The one presented here has the following novel characteristics.

1. Successive combinations differ by a prefix shift. There is no other algorithm for generating combinations with this feature. In some applications combinations are represented in a single computer word; our algorithm is very fast in this scenario. It is also very suitable for hardware implementation.
2. Successive combinations differ by one or two transpositions of a 0 and a 1. There are other algorithms where successive combinations differ by a single transposition (Tang and Liu [10]). Furthermore, that transposition can be further restricted in various ways. For example, so that only 0s are between the transposed bits (Eades and McKay [5]), or such that the transposed bits are adjacent or have only one bit between (Chase [2]). Along with ours, these other variants are ably discussed in Knuth [8].
3. The list is circular; the first and last bitstrings differ by a prefix shift.
4. The list for (s, t) begins with the list for $(s - 1, t)$. Usually, this property is incompatible with Property 3, relative to the elementary operation used to transform one string to the next. For example, colex order has Property 4 but not Property 3. Furthermore, when the elements are expressed as $c_1 c_2 \dots c_t \in \mathbf{C}(s, t)$, the list has the genlex property.
5. The algorithm can be implemented so that *in the worst case* only a small number of operations are done between successive combinations, *independent* of s and t . Such algorithms are said to be loopless, an expression coined by Ehrlich [6]. The algorithm can also be implemented to be loopless and branchless (no *if*-statements).
6. Unlike other Gray codes for combinations, this one has a simple ranking function whose running time is $O(n)$ arithmetic operations.
7. Unlike every other recursive Gray code definition for combinations, (5) has the remarkable property that it involves no reversal of lists.
8. The list is remarkably similar to the colex list for combinations.

The listing discussed here appears in Knuth's preface [8]. The output of the algorithm is illustrated in Figure 26 on page 17. He refers to the listing as suffix-rotated (since he indexes the bitstrings $b_{n-1} \dots b_1 b_0$). See also Exercise 55 on page 30 and its solution on page 97.

To overview the remainder of the paper, Section 2 gives several definitions of cool-lex and proves that they are equivalent, Section 3 provides algorithms and implementations, Section 4 contains the ranking function for cool-lex, Section 5 discusses the genlex property, and Section 6 concludes with several open problems and an extension to permutations of a multi-set.

2 Cool-lex Definitions

In this section, we provide one iterative definition, and two recursive definitions for cool-lex. Theorem 1 proves that all three definitions are equivalent, and gives several immediate consequences. We also provide an iterative and recursive definition for colex.

2.1 Preliminaries and Notation

Before defining the cool-lex order, we introduce a number of secondary definitions. Let $\mathbf{S} = \mathbf{s}_1, \mathbf{s}_2, \dots, \mathbf{s}_m$ be a sequence of strings, let \mathbf{b} , \mathbf{c} , and \mathbf{d} be individual strings, let x be a symbol, let $k \geq 0$, and let $1 \leq i \leq m$. The string \mathbf{bc} is obtained by appending \mathbf{c} to the end of \mathbf{b} . If $\mathbf{d} = \mathbf{bc}$, then \mathbf{b} is a *prefix* of \mathbf{d} , and \mathbf{c} is a *suffix* of \mathbf{d} . \mathbf{Sb} is the sequence of strings $\mathbf{s}_1\mathbf{b}, \mathbf{s}_2\mathbf{b}, \dots, \mathbf{s}_m\mathbf{b}$. Also, x^k is the string with symbol x repeated k times. Let $\mathbf{S}[i] = \mathbf{s}_i$. We frequently access the first and last strings in a sequence, so if \mathbf{S} is non-empty, then $\text{first}(\mathbf{S}) = \mathbf{s}_1$ and $\text{last}(\mathbf{S}) = \mathbf{s}_m$. If \mathbf{S} contains at least two strings, then $\overrightarrow{\mathbf{S}}$ is the rotated sequence of strings $\mathbf{s}_2, \mathbf{s}_3, \dots, \mathbf{s}_m, \mathbf{s}_1$; otherwise if \mathbf{S} does not contain at least two strings, then $\overrightarrow{\mathbf{S}} = \mathbf{S}$. In this paper, every string will be binary, and every symbol will be in $\{0, 1\}$.

When \mathbf{b} is a bitstring of length n , let $l(\mathbf{b})$ be the length of its shortest prefix ending in 010 or 011, or n if no such prefix exists. Let $p(\mathbf{b})$ be the prefix of \mathbf{b} that has length $l(\mathbf{b})$, and let $s(\mathbf{b})$ be the suffix such that $\mathbf{b} = p(\mathbf{b})s(\mathbf{b})$. Let $\sigma(\mathbf{b})$ be the result of rotating $p(\mathbf{b})$ by one position to the right, and appending $s(\mathbf{b})$. Recursively define $\sigma^i(\mathbf{b}) = \sigma(\sigma^{i-1}(\mathbf{b}))$, where $\sigma^0(\mathbf{b}) = \mathbf{b}$.

Properties of σ The strings 1^t0^s and $1^{t-1}0^s1$ play special roles in cool-lex, and their importance and relationship are given by the following three remarks. Remark 4 and Lemma 1 show how transpositions can take the place of rotations.

Remark 1. $\sigma(\mathbf{b})0 = \sigma(\mathbf{b}0)$ if and only if $\mathbf{b} \neq 1^{t-1}0^s1$.

Remark 2. $\sigma(\mathbf{b})1 = \sigma(\mathbf{b}1)$ if and only if $\mathbf{b} \neq 1^t0^s$.

Remark 3. $\sigma(1^{t-1}0^s1) = 1^t0^s$.

Remark 4. $\sigma(\mathbf{b}) = \sigma(p(\mathbf{b}))s(\mathbf{b})$.

Lemma 1. $\sigma(\mathbf{b})$ can be obtained from \mathbf{b} by transposing one or two pairs of bits.

Proof. If $p(\mathbf{b})$ does not end in 010 or 011, then $\mathbf{b} = 1^t 0^s$ and $\sigma(\mathbf{b}) = 01^t 0^{s-1}$, or $\mathbf{b} = 1^{t-1} 0^s 1$ and $\sigma(\mathbf{b}) = 1^t 0^s$. In both of these cases, $\sigma(\mathbf{b})$ can be obtained from \mathbf{b} with one transposition. Otherwise, $p(\mathbf{b})$ does end in 010 or 011 so it must be of the form $00^i 10$, $11^i 00^j 10$, $00^i 11$, or $11^i 00^j 11$, where $i, j \geq 0$. For each case we verify the claim by illustrating the first transposed positions in $p(\mathbf{b})$ using underlines, and if necessary, the second transposed positions in $p(\mathbf{b})$ using overlines. Remark 4 justifies why the transpositions are contained within $p(\mathbf{b})$.

Case 1: $\sigma(00^i 10) = 00^i \underline{10} = 00^i 01$.

Case 2: $\sigma(11^i 00^j 10) = \underline{11^i} \underline{00^j} \overline{10} = 01^i 10^j \overline{10} = 011^i 00^j 1$.

Case 3: $\sigma(00^i 11) = \underline{00^i} \underline{11} = 100^i 1$.

Case 4: $\sigma(11^i 00^j 11) = 11^i \underline{00^j} \underline{11} = 111^i 00^j 1$. \square

2.2 Iterative Definition

Formally, the iterative definition of cool-lex with s zeros and t ones is

$$\mathbf{R}_{s,t} = \sigma^0(\mathbf{b}), \sigma^1(\mathbf{b}), \sigma^2(\mathbf{b}), \dots, \sigma^z(\mathbf{b}) \quad (1)$$

where $\mathbf{b} = 1^t 0^s$ and $z = \binom{s+t}{t} - 1$. When $s = 1$ or $t = 1$, the strings in $\mathbf{R}_{s,t}$ are given explicitly by the following two remarks. The center column of Figure 1 gives $\mathbf{R}_{3,3}$.

Remark 5. $\mathbf{R}_{1,t} = 1^t 0, 01^t, 101^{t-1}, 1^2 01^{t-2}, \dots, 1^{t-1} 01$.

Remark 6. $\mathbf{R}_{s,1} = 10^s, 010^{s-1}, 0^2 10^{s-2}, \dots, 0^s 1$.

To complement the iterative definition of cool-lex, let us consider the well-known iterative definition of colex [8], the lexicographic order applied to the reversal of strings, which begins with $1^t 0^s y$ and ends with $0^s 1^t$. Colex has many uses, for example in Frankl's now standard proof of the Kruskal-Katona Theorem [11]. Let \mathbf{b} be a bitstring of length n . Given that $\mathbf{b} \neq 0^s 1^t$, let $l'(\mathbf{b})$ be the length of the shortest prefix in \mathbf{b} that ends in 10, let $p'(\mathbf{b})$ be the prefix of \mathbf{b} that has length $l'(\mathbf{b})$, and let $s'(\mathbf{b})$ be the suffix of \mathbf{b} such that $\mathbf{b} = p'(\mathbf{b})s'(\mathbf{b})$. Let $\varsigma(\mathbf{b})$ be the result of replacing $p'(\mathbf{b}) = 0^i 1^j 0$ by $1^{j-1} 0^{i+1} 1$ and appending $s'(\mathbf{b})$. Recursively define $\varsigma^i(\mathbf{b}) = \varsigma(\varsigma^{i-1}(\mathbf{b}))$, where $\varsigma^0(\mathbf{b}) = \mathbf{b}$. Notice that $\varsigma(\mathbf{b})$ is well-defined, except for $\mathbf{b} = 0^s 1^t$, which is the last string in colex. The iterative definition of colex with s zeros and t ones is

$$\mathbf{I}_{s,t} = \varsigma^0(\mathbf{b}), \varsigma^1(\mathbf{b}), \varsigma^2(\mathbf{b}), \dots, \varsigma^z(\mathbf{b}) \quad (2)$$

where $\mathbf{b} = 1^t 0^s$ and $z = \binom{s+t}{t} - 1$. When $s = 1$ or $t = 1$, the strings in $\mathbf{I}_{s,t}$ are given explicitly by the following two remarks. The third column of Figure 2 gives $\mathbf{I}_{3,3}$.

Remark 7. $\mathbf{I}_{1,t} = 1^t 0, 1^{t-1} 01, 1^{t-2} 01^2, 1^{t-3} 01^3, \dots, 01^t$.

Remark 8. $\mathbf{I}_{s,1} = 10^s, 010^{s-1}, 0^2 10^{s-2}, \dots, 0^s 1$.

$\mathbf{M}_{3,2}$	$\overrightarrow{\mathbf{M}_{3,2}}$	$\mathbf{M}_{2,3}$	$\mathbf{M}_{3,3} = \mathbf{R}_{3,3}$	$l(\mathbf{b})$	$p(\mathbf{b}) \cdot s(\mathbf{b})$	$\sigma(\mathbf{b})$
		11100	111000	6	111000 ·	011100
		01110	011100	3	011 · 100	101100
		10110	101100	4	1011 · 00	110100
		11010	110100	5	11010 · 0	011010
		01101	011010	3	011 · 010	101010
		10101	101010	4	1010 · 10	010110
		01011	010110	3	010 · 110	001110
		00111	001110	4	0011 · 10	100110
		10011	100110	5	10011 · 0	110010
		11001	110010	6	110010 ·	011001
11000	01100		011001	3	011 · 001	101001
01100	10100		101001	4	1010 · 01	010101
10100	01010		010101	3	010 · 101	001101
01010	00110		001101	4	0011 · 01	100101
00110	10010		100101	5	10010 · 1	010011
10010	01001		010011	3	010 · 011	001011
01001	00101		001011	4	0010 · 11	000111
00101	00011		000111	5	00011 · 1	100011
00011	10001		100011	6	100011 ·	110001
10001	11000		110001	6	110001 ·	111000

Fig. 1. Recursive and iterative structure of cool-lex with $\mathbf{M}_{3,3} = \mathbf{R}_{3,3}$ in the middle column. The leftmost three columns show its recursive structure since $\mathbf{M}_{3,3} = \mathbf{M}_{2,3}0, \overrightarrow{\mathbf{M}_{3,2}1}$. The rightmost three columns show its iterative structure since each string, \mathbf{b} in $\mathbf{R}_{3,3}$, is broken into its prefix $p(\mathbf{b})$ of length $l(\mathbf{b})$, and its suffix $s(\mathbf{b})$. The prefix is rotated by one position to the right to obtain $\sigma(\mathbf{b})$, which is the next string in $\mathbf{R}_{3,3}$.

$\mathbf{L}_{3,2}$	$\mathbf{L}_{2,3}$	$\mathbf{L}_{3,3} = \mathbf{I}_{3,3}$	$l'(\mathbf{b})$	$p'(\mathbf{b}) \cdot s'(\mathbf{b})$	$\varsigma(\mathbf{b})$
	11100	111000	4	1110 · 00	110100
	11010	110100	3	110 · 100	101100
	10110	101100	2	10 · 1100	011100
	01110	011100	5	01110 · 0	110010
	11001	110010	3	110 · 010	101010
	10101	101010	2	10 · 1010	011010
	01101	011010	4	0110 · 10	100110
	10011	100110	2	10 · 0110	010110
	01011	010110	3	010 · 110	001110
	00111	001110	6	001110 ·	110001
11000		110001	3	110 · 001	101001
10100		101001	2	10 · 1001	011001
01100		011001	4	0110 · 01	100101
10010		100101	2	10 · 0101	010101
01010		010101	3	010 · 101	001101
00110		001101	5	00110 · 1	100011
10001		100011	2	10 · 0011	010011
01001		010011	3	010 · 011	001011
00101		001011	4	0010 · 11	000111
00011		000111	-	-	-

Fig. 2. Recursive and iterative structure of colex with $\mathbf{L}_{3,3} = \mathbf{I}_{3,3}$ in the third column. The leftmost two columns show its recursive structure since $\mathbf{L}_{3,3} = \mathbf{L}_{2,3}0, \mathbf{L}_{3,2}1$. The rightmost three columns show its iterative structure since each string, \mathbf{b} in $\mathbf{I}_{3,3}$, is broken into its prefix $p'(\mathbf{b})$ of length $l'(\mathbf{b})$, and its suffix $s'(\mathbf{b})$. The prefix is updated to obtain $\varsigma(\mathbf{b})$, which is the next string in $\mathbf{I}_{3,3}$.

2.3 Recursive Definitions

Although we presented an iterative definition of colex, it is perhaps more commonly expressed recursively. We use $\mathbf{L}_{s,t}$ in the recursive definition, and we note that $\mathbf{I}_{s,t} = \mathbf{L}_{s,t}$ [8]. The colex list \mathbf{L}_{st} is given by the following:

$$\mathbf{L}_{st} = \mathbf{L}_{(s-1)t}0, \mathbf{L}_{s(t-1)}1 \quad (3)$$

where $\mathbf{L}_{0t} = 1^t$ and $\mathbf{L}_{s0} = 0^s$. Interestingly, cool-lex can be defined in a very similar manner. The cool-lex list \mathbf{M}_{st} is given by the following:

$$\mathbf{M}_{st} = \mathbf{M}_{(s-1)t}0, \overrightarrow{\mathbf{M}_{s(t-1)}}1 \quad (4)$$

where $\mathbf{M}_{0,t} = 1^t$ and $\mathbf{M}_{s,0} = 0^s$. Equations (3) and (4) imply that colex can be transformed into cool-lex by a series of sublist manipulations. Figure 3 illustrates this transformation when $s = t = 3$. Remark 9 follows immediately from (4).

Remark 9. Bitstrings with s zeros and t ones appear exactly once in $\mathbf{M}_{s,t}$.

111000	111000	111000
<u>110</u> 100	<u>110</u> 100	011100
<u>101</u> 100	011100	101100
<u>011</u> 100	<u>101</u> 100	110100
<u>1100</u> 10	<u>1100</u> 10	011010
<u>1010</u> 10	011010	101010
<u>0110</u> 10	101010	010110
<u>1001</u> 10	010110	001110
<u>0101</u> 10	001110	100110
<u>0011</u> 10	<u>1001</u> 10	110010
<u>11000</u> 1	<u>11000</u> 1	011001
<u>10100</u> 1	011001	101001
<u>01100</u> 1	101001	010101
<u>10010</u> 1	010101	001101
<u>01010</u> 1	001101	100101
<u>00110</u> 1	100101	010011
<u>10001</u> 1	010011	001011
<u>01001</u> 1	001011	000111
<u>00101</u> 1	000111	100011
<u>00011</u> 1	<u>10001</u> 1	110001

Fig. 3. From left to right: transforming colex into cool-lex. In each rectangle, the substrings have a common suffix beginning in 1 and are cyclically moved up one row.

Although the representation of cool-lex in (4) has certain advantages, it can also be useful to have a recursive definition that does not use $\overrightarrow{\mathbf{S}}$. By using the same base cases, we can define cool-lex recursively by $\mathbf{W}'_{st} = 1^t 0^s$, \mathbf{W}_{st} where

$$\mathbf{W}_{st} = \mathbf{W}_{(s-1)t}0, \mathbf{W}_{s(t-1)}1, 1^{t-1}0^s1 \quad (5)$$

In fact, this definition is used in [8] and in a conference paper containing preliminary results [12]. When $s = 1$ or $t = 1$, the strings in $\mathbf{M}_{s,t}$ and $\mathbf{W}'_{s,t}$ are given explicitly in the following two remarks. Figure 1 gives $\mathbf{W}'_{3,3}$.

Remark 10. $\mathbf{M}_{1,t} = \mathbf{W}'_{1,t} = 1^t 0, 01^t, 101^{t-1}, 1^2 01^{t-2}, \dots, 1^{t-1} 01$.

Remark 11. $\mathbf{M}_{s,1} = \mathbf{W}'_{s,1} = 10^s, 010^{s-1}, 0^2 10^{s-2}, \dots, 0^s 1$.

Lemma 3 proves that $\mathbf{M}_{s,t} = \mathbf{W}'_{s,t}$. One advantage of $\mathbf{W}'_{s,t}$ is that it is easy to identify the first and last strings in the cool-lex. We will also find it useful to know the second string in cool-lex order, which we compute using $\mathbf{M}_{s,t}$.

Lemma 2. (a) $first(\mathbf{W}'_{s,t}) = 1^t 0^s$
(b) $last(\mathbf{W}'_{s,t}) = 1^{t-1} 0^s 1$
(c) $second(\mathbf{M}_{s,t}) = 01^t 0^{s-1}$ for $s, t > 1$

Proof. Parts (a) and (b) follow immediately from the definitions. For part (c),

$$second(\mathbf{M}_{s,t}) = second(\mathbf{M}_{s-1,t})0 = \dots = second(\mathbf{M}_{1,t})0^{s-1} = 01^t 0^{s-1}$$

□

2.4 Equivalence of Definitions

Now we are ready to prove the main result of this section, Theorem 1. We first prove that the two recursive definitions of cool-lex, $\mathbf{M}_{s,t}$ and $\mathbf{W}'_{s,t}$ are equivalent, and then we prove that these definitions are equivalent to the iterative definition of cool-lex, $\mathbf{R}_{s,t}$.

Lemma 3. $\mathbf{M}_{s,t} = \mathbf{W}'_{s,t}$.

Proof. From Remarks 10 and 11, the result is true when $s = 1$ or $t = 1$. Otherwise, suppose that $s, t > 1$ and inductively assume that $\mathbf{M}_{i,j} = \mathbf{W}'_{i,j}$ whenever $i < s$ and $j < t$. Then we have the following:

$$\begin{aligned} \mathbf{W}'_{s,t} &= 1^t 0^{s-1} 0, \mathbf{W}_{s,t} \\ &= 1^t 0^{s-1} 0, \mathbf{W}_{s-1,t} 0, \mathbf{W}_{s,t-1} 1, 1^{t-1} 0^s 1 \\ &= \mathbf{W}'_{s-1,t} 0, \mathbf{W}_{s,t-1} 1, 1^{t-1} 0^s 1 \\ &= \mathbf{M}_{s-1,t} 0, \mathbf{W}_{s,t-1} 1, 1^{t-1} 0^s 1 \\ &= \mathbf{M}_{s-1,t} 0, \overrightarrow{(1^{t-1} 0^s 1, \mathbf{W}_{s,t-1} 1)} \\ &= \mathbf{M}_{s-1,t} 0, \overrightarrow{(\mathbf{W}'_{s,t-1} 1)} \\ &= \mathbf{M}_{s-1,t} 0, \overrightarrow{(\mathbf{M}_{s,t-1} 1)} \\ &= \mathbf{M}_{s,t} \end{aligned}$$

□

Lemma 4. $\mathbf{M}_{s,t} = \mathbf{R}_{s,t}$.

Proof. Remarks 5, 6, 10, and 11 provide the result when $s = 1$ or $t = 1$. Otherwise, suppose that $s, t > 1$ and inductively assume that $\mathbf{M}_{i,j} = \mathbf{W}_{i,j}$ whenever $i < s$ and $j < t$. The following list gives an overview of $\mathbf{M}_{s,t} = \mathbf{M}_{s-1,t}0, \overrightarrow{\mathbf{M}_{s,t-1}1}$, with a horizontal line separating the two sublists. We wish to show that each successive string in $\mathbf{M}_{s,t}$ is the result of applying σ to the previous string.

$$\begin{array}{l} 11^{t-2}100^{s-2}0 \\ 011^{t-2}10^{s-2}0 \end{array} \quad (\text{S1})$$

$$\begin{array}{l} \vdots \\ 11^{t-2}00^{s-2}10 \\ \hline 011^{t-2}00^{s-2}1 \end{array} \quad \begin{array}{l} (\text{S2}) \\ (\text{S3}) \end{array}$$

$$\begin{array}{l} \vdots \\ 1^{t-2}00^{s-2}011 \\ 11^{t-2}00^{s-2}01 \end{array} \quad \begin{array}{l} (\text{S4}) \\ (\text{S5}) \end{array}$$

The strings from (S1) to (S2) are the strings in $\mathbf{M}_{s-1,t}0$. From Remarks 1, 9 and Lemma 2 (b), appending zero does not affect the operation of σ , except for the string labelled (S2). Therefore, the fact that each successive string from (S1) to (S2) is obtained from applying σ is a result of the inductive assumption that $\mathbf{M}_{s-1,t} = \mathbf{R}_{s-1,t}$. Next, notice that applying σ to the string (S2) results in the string (S3).

The strings from (S3) to (S5) are the strings in $\overrightarrow{\mathbf{M}_{s,t-1}1}$. From Remarks 2 and 9 and Lemma 2 (a), appending one does not affect the operation of σ , for every string from (S3) to (S4). Therefore, the fact that each successive string from (S3) to (S4) is obtained from applying σ is a result of the inductive assumption that $\mathbf{M}_{s,t-1} = \mathbf{R}_{s,t-1}$. Finally, notice that applying σ to the string (S4) results in the string (S5). \square

Theorem 1. $\mathbf{R}_{st} = \mathbf{W}'_{st} = \mathbf{M}_{st}$. Moreover,

- The lists are circular.
- The lists contains each bitstring exactly once.
- Successive bitstrings differ by a prefix shift of one position to the right.
- Successive bitstrings differ by the transposition of one or two pairs of bits.
- The first bitstring is 1^t0^s , and the last bitstring is $1^{t-1}0^s1$.

Proof. The first point follows from Remark 3. The second point follows from Remark 9. The third point follows from the definitions of $\mathbf{R}_{s,t}$ and σ . The fourth point follows from Remark 1. The last point follows from Remark 2. \square

3 Algorithms and Implementation

In this section, we concentrate on efficient algorithms for generating cool-lex. In particular, we provide a recursive algorithm, a loopless iterative algorithm, and

a loopless and branchless iterative algorithm, each of which is implemented in a procedural language. We also provide a second loopless and branchless iterative algorithm which is implemented in machine language, and is due to Knuth [8].

3.1 Recursive Algorithm

To generate cool-lex recursively we use the definitions of $\mathbf{W}'_{s,t}$ and $\mathbf{W}_{s,t}$ which we recall here:

$$\begin{aligned}\mathbf{W}'_{st} &= 1^t 0^s, \mathbf{W}_{st} \\ \mathbf{W}_{st} &= \mathbf{W}_{(s-1)t} 0, \mathbf{W}_{s(t-1)} 1, 1^{t-1} 0^s 1\end{aligned}$$

Figure 4 shows the strings in $\mathbf{W}_{s,t}$, where the two long horizontal lines represent the transitions between $\mathbf{W}_{(s-1)t} 0$, $\mathbf{W}_{s(t-1)} 1$, and $1^{t-1} 0^s 1$. The left column shows the base case of $t = 1$, the middle column shows the base case of $s = 1$, and the right column shows the remaining case of $s, t > 1$. The short underlines, and overlines, represent which bits are transposed at each interface.

$010^{s-2}0$	$\overline{011^{t-2}1}$	$011^{t-2}10^{s-2}0$
\vdots	\vdots	\vdots
\vdots	\vdots	$\underline{11^{t-2}00^{s-2}\overline{10}}$
\vdots	\vdots	$01^{t-2}10^{s-2}01$
\vdots	\vdots	\vdots
$\underline{0^{s-2}0\underline{10}}$	$\underline{1^{t-2}\underline{0}\underline{11}}$	$\underline{1^{t-2}\underline{000^{s-2}}\underline{11}}$
$\overline{00^{s-2}01}$	$\overline{11^{t-2}01}$	$\overline{1^{t-2}100^{s-2}01}$

Fig. 4. Illustrating the transpositions at the two interfaces in $\mathbf{W}_{s,t}$.

In the right column, the transposed bits at the first interface are at positions $(1, t)$ and $(n - 1, n)$, and at the second interface are at positions $(t - 1, n - 1)$ (Lemma 1). To generate all of the strings in $\mathbf{W}'_{s,t}$ we initialize \mathbf{b} to $1^t 0^s$, visit this string, and then recursively generate $\mathbf{W}_{s,t}$ using the algorithm below. The array \mathbf{b} is indexed starting at position 1, and we assume $s, t > 0$. The initial call is `swap(1, t+1); visit(b); gen(s, t);`. Since every recursive call is followed by a visit, the algorithm runs in constant amortized time.

```
static void gen ( int s, int t ) {
    if ( s > 1 ) { gen( s-1, t );
        swap( 1, t ); swap( s+t, s+t-1 ); visit( b ); }
    if ( t > 1 ) { gen( s, t-1 );
        swap( t-1, s+t-1 ); visit( b ); }
}
```

3.2 Iterative Algorithms

Loopless Algorithm We now present the iterative loopless implementation. Again we begin by initializing \mathbf{b} to $1^t 0^s$ and by visiting this string. In this case the array indexing is 0 based. It is useful to maintain a variable x , which is the smallest index for which $\mathbf{b}[x-1] == 0$ and $\mathbf{b}[x] == 1$. In terms of prefix shifts, the code to obtain the next bitstring and to update x is amazingly simple.

```
shift( b, ++x );
if (b[0] == 0 && b[1] == 1) x = 1;
```

To generate the next bitstring by transpositions it is useful to maintain another variable y , which is the smallest index for which $\mathbf{b}[y] == 0$. Referring back to Figure 4 we observe that in every case $\mathbf{b}[x]$ becomes 0 and $\mathbf{b}[y]$ becomes 1. The test $\mathbf{b}[x+1] == 0$ determines whether we are at the first or the second interface. If we are at the first interface, then we set $\mathbf{b}[x+1]$ to 1 and $\mathbf{b}[0]$ to 0. It now remains to update x and y . At the second interface they are simply incremented. At the first interface y always becomes 0; also, x is incremented unless $y = 0$, in which case x becomes 1 (see Remark 11).

```
static void iterate ( int s, int t ) {
    b[t] = 1; b[0] = 0;
    visit( b );
    int x = 1, y = 0;
    while (x < n-1) {
        b[x++] = 0; b[y++] = 1;    /* X(s,t) */
        if (b[x] == 0) {
            b[x] = 1; b[0] = 0;    /* Y(s,t) */
            if (y > 1) x = 1;      /* Z(s,t) */
            y = 0; }
        visit( b ); } }
```

The structure of the implementation allows us to completely determine the number of times each statement in the code is executed. Call the relevant quantities $X(s, t)$, $Y(s, t)$, and $Z(s, t)$ corresponding to the various statements as shown above. I.e., $Y(s, t)$ is the number of times $\mathbf{b}[x] == 0$ is true and $Z(s, t)$ is the number of times $y > 1$ is true. We find that

$$X(s, t) = \binom{s+t}{t} - 1, \quad Y(s, t) = \binom{s+t-1}{t}, \quad Z(s, t) = \binom{s+t-2}{t-1}.$$

Loopless and Branchless Algorithm The previous loopless algorithm generates cool-lex by transposing either one pair or two pairs of bits at each step. Interestingly, cool-lex can also be generated by *always* swapping two pairs of bits. In particular, by maintaining the variables as before, each successive string can be obtained by `swap(x, y); swap(0, x+1);`. As was previously done, σ will be manually applied when $\mathbf{b} = \text{first}(\mathbf{R}_{s,t}) = 1^t 0^s$ and will not be applied when $\mathbf{b} = \text{last}(\mathbf{R}_{s,t}) = 1^{t-1} 0^s 1$. In all other cases, $p(\mathbf{b})$ will end in 010 or 011, which explains the restriction placed on the following lemma.

Lemma 5. *If $p(\mathbf{b})$ ends in 010 or 011, then $\sigma(\mathbf{b})$ can be obtained from \mathbf{b} by transposing bits (x, y) , followed by transposing bits $(0, x + 1)$.*

Proof. Since $p(\mathbf{b})$ ends in 010 or 011, then it must be of the form 00^i10 , 11^i00^j10 , 00^i11 , or 11^i00^j11 , where $a, b \geq 0$. By Remark 4 we need only transpose bits in $p(\mathbf{b})$, and for each case we verify the claim by illustrating the transposition of positions (x, y) using underlines, and the transposition of positions $(0, x + 1)$ using overlines.

$$\text{Case 1: } \sigma(00^i10) = \overline{00^i}\underline{10} = \overline{10^i}0\overline{0} = 00^i01.$$

$$\text{Case 2: } \sigma(11^i00^j10) = \overline{11^i}\underline{00^j}\underline{10} = \overline{11^i}10^j0\overline{0} = 011^i00^j1.$$

$$\text{Case 3: } \sigma(00^i11) = \overline{00^i}\underline{11} = \overline{10^i}0\overline{1} = 100^i1.$$

$$\text{Case 4: } \sigma(11^i00^j11) = \overline{11^i}\underline{00^j}\underline{11} = \overline{11^i}10^j0\overline{1} = 111^i00^j1. \quad \square$$

Given the correct values of x and y , Lemma 5 allows us to generate the next string without branching. Once the next string has been generated we can easily compute the correct values of x and y . In particular, the value of y is incremented by one, unless the first bit is set to 0, in which case set $y = 0$. Likewise, the value of x is incremented by one, unless the first two bits are set to 01, in which case set $x = 1$. This observation allows us to replace the loop in the algorithm from Section 3.2 with the following statements (the looping condition must also change from `while (x < n-1)` to `while (x < n)` since x is now updated at the end of the loop).

```

b[x] = 0;
b[y] = 1;
b[0] = b[x + 1];
b[x + 1] = 1;
x = x * (1 - b[1] + b[1] * b[0]) + 1;
y = (y + 1) * b[0];
visit( b );

```

Although the resulting algorithm is no faster than the previous loopless algorithm, it does have the interesting property that it is both loopless and branchless. In particular, each successive strings is generated using only the variables x and y , and the same six assignment statements.

3.3 Implementation in Computer Words

The final implementation we present is of a different nature than the previous three. In this case we assume that our n -bit binary string can fit inside of a single machine word, and we operate on this word using machine language. By using shifts, bitmasks, and arithmetic, there are a number of ways to accomplish this goal. The approach we follow here is due to Knuth [8], and it gives a loopless

and branchless MMIX implementation. To understand the algorithm, we need the following two lemmas, which show how the operation of σ can be simulated by using addition and subtraction on words. Again, we focus only on $p(\mathbf{b})$ due to Remark 4

Lemma 6. *If $p(\mathbf{b}) = 1^x 00^y 10$, then $\sigma(\mathbf{b}) = \mathbf{b} + \mathbf{c}$, for $\mathbf{c} = 1^x 00^y 1$.*

Proof. To verify that we can obtain $\sigma(1^x 00^y 10) = 01^x 0^y 01$ by adding \mathbf{c} , we write each string from right to left:

$$\begin{array}{r} 010^y 01^x \\ + 010^y 01^x \\ \hline 100^y 1^x 0 \end{array}$$

□

Lemma 7. *If $p(\mathbf{b}) = 1^x 00^y 11$, then $\sigma(\mathbf{b}) = \mathbf{b} - \mathbf{c}$, for $\mathbf{c} = 0^x 11^y$.*

Proof. To verify that we can obtain $\sigma(1^x 00^y 11) = 11^x 0^y 01$ by subtracting \mathbf{c} , we write each string from right to left:

$$\begin{array}{r} 110^y 01^x \\ - 001^y 10^x \\ \hline 100^y 11^x \end{array}$$

□

For the implementation, we assume that every register has length w , and that $n = s + t < w$. Registers $R0$ and $R1$ are used as temporary variables, and the binary strings will be generated in $R2$, which is initialized by setting $R2 = 1^t 0^{w-t}$ (this can be done by $R2 \leftarrow (1 \ll t) - 1$). The operation $_$ is a specialized form of subtraction, where the result of $i _ j$ is $i - j$ if $j \leq i$, and is 0 if $j > i$. Although this operation is not available in all machine languages, it is available in MMIX, and can easily be simulated using other instructions. The algorithm runs until the $(n + 1)$ st bit is set to 1.

```
R0 <- R2 & (R2+1)
R1 <- R0 xor (R0-1)
R0 <- R1+1
R1 <- R1 & R2
R0 <- (R0 & R2) _ 1
R2 <- R2 + R1 - R0
```

To understand the implementation, suppose $R2 = 1^x 00^y 1 d s(\mathbf{b})$ where d is a single bit. The first statement places $0^{x+y+1} 1 d s(\mathbf{b})$ into $R0$ (this is the value of $R2$ with leading 1s changed to 0s). The second statement places $1^{x+y+2} 0^{w-x-y-2}$ into $R1$ (this is a mask for the shortest prefix ending 01 in $R2$). The third statement places $0^{x+y+2} d 0^{w-x-y-3}$ into $R1$. The fourth statement puts the value of $1^x 0^y 01 0^{w-x-y-2}$ into $R1$ (this is the shortest prefix in $R2$ ending in 01, with the remaining bits set to 0). If $d = 0$, then the fifth statement puts the value

of 0^w into $R0$, and then the sixth statement puts the correct value into $R2$ via Lemma 6. If $d = 1$, then the fifth statement puts the value of $1^{x+y+2}0^{w-x-y-2}$ into $R0$, and then the sixth statement puts the correct value of into $R2$ via Lemma 7 since $R0 - R1 = 0^x 1^y 0^{w-x-y}$.

4 Ranking Algorithm

In this section we examine the ranking functions of colex and cool-lex, and this provides another interesting link between the two lists. Given a listing of combinatorial structures, the *rank* of a particular structure is the number of structures that precede it in the listing.

Given an (s, t) -combination represented as a bitstring $b_1 b_2 \dots b_n$ the corresponding set elements can be listed as $c_1 < c_2 < \dots < c_t$ where c_i is the position of the i -th 1 in the bitstring. As is well-known ([8],[11]) in colex order the rank of $c_1 c_2 \dots c_t$ is

$$\sum_{j=1}^t \binom{c_j - 1}{j}. \quad (6)$$

As we see in the statement of the theorem below, in cool-lex order there is a very similar rank function. Let $rank(c_1 c_2 \dots c_t)$ denote the rank of $c_1 c_2 \dots c_t \in \mathbf{C}(s, t)$ in our order.

Theorem 2. *Let r be the smallest index such that $c_r > r$ (so that $c_{r-1} = r - 1$).*

$$rank(c_1 c_2 \dots c_t) = \binom{c_r}{r} - 1 + \sum_{j=r+1}^t \left(\binom{c_j - 1}{j} - 1 \right), \quad (7)$$

Proof. Directly from the recursive construction (5) we have

$$rank(b_1 b_2 \dots b_n) = \begin{cases} rank(b_1 b_2 \dots b_{n-1}) & \text{if } b_n = 0, \\ \binom{s+t}{t} - 1 & \text{if } b_1 b_2 \dots b_n = 1^{t-1} 0^s 1, \\ \binom{s+t-1}{t-1} - 1 + rank(b_1 b_2 \dots b_{n-1}) & \text{otherwise.} \end{cases}$$

Let us now consider the rank in terms of the corresponding list of elements $1 \leq c_1 < c_2 < \dots < c_t$. The case $rank(b_1 b_2 \dots b_n) = rank(b_1 b_2 \dots b_{n-1}) = rank(b_1 b_2 \dots b_{n-2})$ will continue to apply until $b_{n-k} = 1$; i.e., until $n - k = c_{t-1}$. Hence the number of 0s and 1s to the left of position c_{t-1} in $b_1 b_2 \dots b_n$ is $c_{t-1} - 1$, which leads us to the expression below.

$$rank(c_1 c_2 \dots c_t) = \begin{cases} \binom{c_t}{t} - 1 & \text{if } c_t = n \text{ and } c_{t-1} = t - 1 \\ \binom{c_{t-1}-1}{t-1} - 1 + rank(c_1 c_2 \dots c_{t-1}) & \text{otherwise.} \end{cases}$$

As in (6) and (7), the recursion above has the remarkable and useful property that it depends only on t and not on s . In other words, the cool-lex lists begin with cool-lex lists with smaller s values (fewer zeros). \square

The ranking function can also be written recursively, as shown below.

$$\text{rank}(c_1 c_t \cdots c_t) = \text{rank}(c_1 c_2 \cdots c_{t-1}) + \binom{c_r - 1}{r - 1} + r - t - 1. \quad (8)$$

Using standard techniques, as explained for example in [8] the expression in (7) can be evaluated in $O(n)$ arithmetic operations.

5 Genlex

A list of strings, $\mathbf{S} = \mathbf{s}_1, \mathbf{s}_2, \dots, \mathbf{s}_m$, is *genlex* [8] if every suffix appears within consecutive strings in \mathbf{S} . That is, \mathbf{S} is genlex unless there exists a string \mathbf{x} , and integers i, j, k with $1 \leq i < j < k \leq m$, such that \mathbf{x} is a suffix of \mathbf{s}_i and \mathbf{s}_k , but is not a suffix of \mathbf{s}_j . This property is common to several other orderings for combinations [8], and depending on the setting may be defined for prefixes instead of suffixes.

For example, when $s = 2$ and $t = 3$, Figure 5 illustrates that colex is genlex when it is represented by binary strings. This implies that colex is also genlex when it is represented by the elements contained in each combination, which is how genlex is defined for combinations in [8]. To verify this fact, notice that every one element suffix that appears in the list (3, 4, and 5) does so in consecutive strings, as does every two element suffix (23, 24, 25, 34, 35, and 45), and every three element suffix appears exactly once since $t = 3$. Remark 12 states that colex actually has a stronger property than genlex. In particular, the strings with a particular suffix give a smaller colex list.

Remark 12. If \mathbf{x} is a binary string with s' zeros and t' ones, then the strings in $\mathbf{L}_{s-s', t-t'} \mathbf{x}$ appear in the same order, and are consecutive, within $\mathbf{L}_{s,t}$. Moreover, no other strings in $\mathbf{L}_{s,t}$ have \mathbf{x} as a suffix.

On the other hand, Figure 5 shows that cool-lex is not genlex when it is represented by binary strings since the suffix 01 appears in three non-consecutive strings. However, we will show that cool-lex is genlex when it is represented by the elements contained in each combination. In terms of binary strings, this is equivalent to showing that every suffix that begins with a 1 appears in consecutive strings in cool-lex. In Theorem 3, we will prove a result that is stronger than genlex for suffixes that begin with 1, and we will prove a result that is weaker than genlex for suffixes that begin with 0. For intuition, the reader may wish to refer to Figure 3, which illustrates the suffix properties that are maintained when colex is transformed into cool-lex. Before stating the theorem, we mention the following remark that follows directly from Remark 9.

Remark 13. If \mathbf{x} is a binary string with s' zeros and t' ones, then the strings in $\mathbf{R}_{s,t}$ with suffix \mathbf{x} are exactly the strings within $\mathbf{R}_{s-s', t-t'} \mathbf{x}$ (or equivalently, within $\overrightarrow{\mathbf{R}_{s-s', t-t'} \mathbf{x}}$).

$\mathbf{R}_{2,3}$		$\mathbf{L}_{2,3}$	
11100	123	11100	123
11010	124	01110	234
10110	134	10110	134
01110	234	11010	124
11001	125	01101	235
10101	135	10101	135
01101	235	01011	245
10011	145	00111	345
01011	245	10011	145
00111	345	11001	125

Fig. 5. For $s = 2$ and $t = 3$, from left to right, colex as bitstrings, colex as elements, cool-lex as bitstrings, and cool-lex as elements. Only the third column is not genlex.

Theorem 3. 1. If $\mathbf{x} = 1\mathbf{x}'$ is a binary string with $s' \geq 0$ zeros and $t' > 0$ ones, then the strings in $\overrightarrow{\mathbf{R}_{s-s', t-t'}\mathbf{x}}$ appear in the same order, and are consecutive, within $\mathbf{R}_{s,t}$.

2. If $\mathbf{x} = 0\mathbf{x}'$ is a binary string with $s' > 0$ zeros and $t' \geq 0$ ones, then the strings in $\overrightarrow{\mathbf{R}_{s-s', t-t'}\mathbf{x}}$, except for $\text{last}(\overrightarrow{\mathbf{R}_{s-s', t-t'}\mathbf{x}})$, appear in the same order, and are consecutive, within $\mathbf{R}_{s,t}$.

Proof. In this proof we will make implicit use of Remark 13, Lemma 2, Remark 9, Theorem 1, and the definition of $\vec{\mathbf{S}}$. For the first result, if $s' \geq s$ or $t' \geq t$, then the result is immediate since $\mathbf{R}_{s-s', t-t'}$ contains at most one string. Otherwise, the string

$$\begin{aligned} 01^{t-t'}0^{s-s'-1}\mathbf{x} &= \text{second}(\mathbf{R}_{s-s', t-t'})\mathbf{x} \\ &= \text{first}(\overrightarrow{\mathbf{R}_{s-s', t-t'}})\mathbf{x} \end{aligned}$$

must occur somewhere within $\mathbf{R}_{s,t}$. Since $\mathbf{x} = 1\mathbf{x}'$, by Lemma 2,

$$\sigma(\overrightarrow{\mathbf{R}_{s-s', t-t'}[i]})\mathbf{x} = \sigma(\overrightarrow{\mathbf{R}_{s-s', t-t'}[i]}\mathbf{x})$$

unless

$$\begin{aligned} \overrightarrow{\mathbf{R}_{s-s', t-t'}[i]} &= 1^{t-t'}0^{s-s'} \\ &= \text{first}(\mathbf{R}_{s-s', t-t'}) \\ &= \text{last}(\overrightarrow{\mathbf{R}_{s-s', t-t'}}). \end{aligned}$$

Therefore, the strings in $\overrightarrow{\mathbf{R}_{s-s', t-t'}\mathbf{x}}$ appear in the same order, and are consecutive, within $\mathbf{R}_{s,t}$.

For the second result, if $s' \geq s$ or $t' \geq t$, then the result is immediate since $\mathbf{R}_{s-s', t-t'}$ contains at most one string. Otherwise, the string

$$\begin{aligned} 01^{t-t'}0^{s-s'-1}\mathbf{x} &= \text{second}(\mathbf{R}_{s-s', t-t'})\mathbf{x} \\ &= \text{first}(\overrightarrow{\mathbf{R}_{s-s', t-t'}})\mathbf{x} \end{aligned}$$

must occur somewhere within $\mathbf{R}_{s,t}$. Since $\mathbf{x} = 0\mathbf{x}'$, by Lemma 1,

$$\sigma(\overrightarrow{\mathbf{R}_{s-s', t-t'}[i]})\mathbf{x} = \sigma(\overrightarrow{\mathbf{R}_{s-s', t-t'}[i]\mathbf{x}})$$

unless

$$\begin{aligned} \overrightarrow{\mathbf{R}_{s-s', t-t'}[i]} &= 1^{t-t'-1}0^{s-s'}1 \\ &= \text{last}(\mathbf{R}_{s-s', t-t'}). \end{aligned}$$

Notice that $\text{last}(\mathbf{R}_{s-s', t-t'})$ is the penultimate string in $\overrightarrow{\mathbf{R}_{s-s', t-t'}}$. Therefore, the strings in $\overrightarrow{\mathbf{R}_{s-s', t-t'}}\mathbf{x}$, except for $\text{last}(\overrightarrow{\mathbf{R}_{s-s', t-t'}})\mathbf{x}$, appear in the same order, and are consecutive, within $\mathbf{R}_{s,t}$. \square

Before concluding this section, we provide a remark that follows directly from the recursive definition of cool-lex (4), and gives a slight strengthening of Theorem 3.

Remark 14. If $\mathbf{x} = 0^{s'}$, then the strings in $\mathbf{R}_{s-s', t-t'}\mathbf{x}$ appear in the same order, and are consecutive, within $\mathbf{R}_{s,t}$.

6 Final Remarks

6.1 Permutations of a Multi-Set

A multi-set is a collection of integers, with repetition allowed. For example, $S = \{2, 1, 1, 0\}$ is the multi-set with one copy of 2, two copies of 1, and one copy of 0. A permutation of a multi-set, or a *multi-perm*, is any arrangement of these integers. Notice that (s, t) -combinations are simply permutations of the multi-set with s copies of 0 and t copies of 1. Several algorithms exist for generating multi-perms [16, 15, 14, 13].

It appears that the iterative definition of cool-lex can be generalized to generate multi-perms. Starting from any multi-perm, let \mathbf{p} be its shortest prefix ending xy where $x < y$. If there is no such prefix, then let \mathbf{p} be the entire multi-perm. If \mathbf{p} is not the entire multi-perm, and if the next symbol after y is z , where $z \leq x$, then let \mathbf{p} instead be the prefix ending in xyz . Rotate the symbols in \mathbf{p} by one position to the right. For example, the permutations of the multi-set $\{2, 1, 1, 0\}$ are:

$$2110, 0211, 2011, 1201, 0121, 1021, 2101, 1210, 1120, 0112, 1012, 1102.$$

Experimental results have validated the effectiveness of this iterative rule in a number of cases, and we hope to prove its correctness in a follow-up paper. When applied to binary multi-sets, the rule reduces to rotating the shortest prefix ending in 010 or 011 (notice that rotating a prefix ending in 011 is equivalent to rotating the same prefix ending in 01).

6.2 Artistic Representations

The iterative cool-lex list $\mathbf{R}_{s,t}$ has been rendered musically by George Tzanetakis and is available for download as a .wav file on the page <http://www.cs.uvic.ca/~ruskey/Publications/Coollex/Coollex.html>. A visual comparison of colex and cool-lex is illustrated artistically in the *The Feast* (ISBN 0-978066-30-98) and is available on <http://www.pmntmrkr.com/>.

6.3 Open Problems

- Is it possible to generate combinations if the allowed operations are further restricted? For example, all permutations can be generated by letting the permutations $(1\ 2)$ and $(1\ 2\ \cdots\ n)$ and their inverses act on the indices, although this is not possible for combinations (for example, try $s = t = 3$).
- What is the fastest combination generator when carefully implemented? It would be interesting to undertake a comparative evaluation in a controlled environment, say of carefully implemented MMIX programs. Testing should be done, in the three cases, depending on whether the combination is represented by a single computer word, an element of $\mathbf{B}(s, t)$, or an element of $\mathbf{C}(s, t)$.
- What is the computational complexity of determining if an arbitrary subset of (s, t) -combinations can be generated by prefix shifts?

References

1. P.F. Corbett, *Rotator Graphs: An Efficient Topology for Point-to-Point Multiprocessor Networks*, IEEE Transactions on Parallel and Distributed Systems, 3 (1992) 622–626
2. P.J. Chase, *Combination Generation and Graylex Ordering*, Congressus Numerantium, 69 (1989) 215–242.
3. P. Diaconis and S. Holmes, *Gray codes for randomization procedures*, Statistical Computing, 4 (1994) 207–302.
4. P. Eades, M. Hickey and R. Read, *Some Hamilton Paths and a Minimal Change Algorithm*, Journal of the ACM, 31 (1984) 19–29.
5. P. Eades and B. McKay, *An Algorithm for Generating Subsets of Fixed Size with a Strong Minimal Change Property*, Information Processing Letters, 19 (1984) 131–133.
6. G. Ehrlich, *Loopless Algorithms for Generating Permutations, Combinations and Other Combinatorial Configurations*, Journal of the ACM, 20 (1973) 500–513.
7. M. Jiang and F. Ruskey, *Determining the Hamilton-connectedness of certain vertex-transitive graphs*, Discrete Mathematics, 133 (1994) 159–170.
8. Donald E. Knuth, *The Art of Computer Programming, Volume 4: Generating all Combinations and Partitions*, Fascicle 3, Addison-Wesley, July 2005, 150 pages.
9. F. Ruskey, *Simple combinatorial Gray codes constructed by reversing sublists*, 4th ISAAC (International Symposium on Algorithms and Computation), Lecture Notes in Computer Science, #762 (1993) 201–208.
10. D.T. Tang and C.N. Liu *Distance-2 Cycle Chaining of Constant Weight Codes*, IEEE Transactions, C-22 (1973) 176–180.

11. D. Stanton and D. White, *Constructive Combinatorics*, Springer-Verlag, 1986.
12. F. Ruskey and A. Williams, *Generating Combinations By Prefix Shifts*, Computing and Combinatorics, 11th Annual International Conference, COCOON 2005, Kunming, China, August 16-29, 2005, Proceedings. Lecture Notes in Computer Science 3595 Springer (2005).
13. V. Vajnovszki, *A loopless algorithm for generating the permutations of a multiset*, Theoretical Computer Science, 307 (2003), 415–431.
14. T. Takaoka, *An $O(1)$ Time Algorithm for Generating Multiset Permutations*, Proceedings of the 10th International Symposium on Algorithms and Computation, 25 (1999), 237–246.
15. J. Korsh and S. Lipschutz, *Generating multiset permutations in constant time*, Journal of Algorithms, 25 (1997), 321–335.
16. P. J. Chase, *Algorithm 383: permutations of a set with repetitions*, CACM, 13 (1970), 368–369.